



# Credit Card Fraud

---

Kevin Gui, Steven Ha




# Intro

---

Our Overarching Idea: Detect Fraudulent Credit Card Transactions

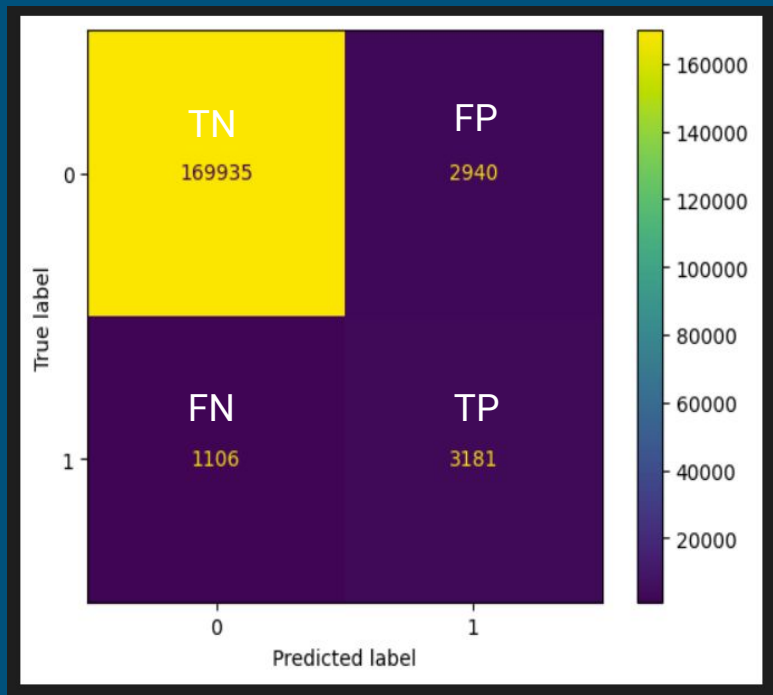
The data is very imbalanced

- Finding the best ML model
- How we improved our model
- Conclusions and Results



# Model Evaluation and Metrics

## Confusion Matrix



Accuracy: 
$$\frac{TP + TN}{TP + TN + FP + FN}$$


Precision: 
$$\frac{TP}{TP + FP}$$

Recall: 
$$\frac{TP}{TP + FN}$$

F1 Score: 
$$2 \times \frac{Precision \times Recall}{Precision + Recall}$$

# AUC/ROC Curve

- Plots TP vs. FP
- Higher AUC score = more accurate model



# Model 1

## Logistic Regression

Accuracy Score : 0.9663415405109448


AUC Score : 0.6107526375093872

Confusion Matrix :

```
[[170884  157]
 [ 5806   315]]
```


Classification Report :

	precision	recall	f1-score	support
False	0.97	1.00	0.98	171041
True	0.67	0.05	0.10	6121



# Models 2/3: Boosting

## AdaBoost



Accuracy Score : 0.970484641175873

AUC Score : 0.8517222221643287


Confusion Matrix :

```
[[170626  415]
 [ 4814 1307]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.97	1.00	0.98	171041
True	0.76	0.21	0.33	6121

## XGBoost



Accuracy Score : 0.9786579514794369

AUC Score : 0.9242500893514191

Confusion Matrix :


```
[[170676  365]
 [ 3416 2705]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	1.00	0.99	171041
True	0.88	0.44	0.59	6121

# Models 4/5: Trees

## Decision Tree Classifier



Accuracy Score : 0.974627741840801

AUC Score : 0.8109088713848963


Confusion Matrix :

```
[[169996  1045]
 [ 3450  2671]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	0.99	0.99	171041
True	0.72	0.44	0.54	6121

## Random Forest



Accuracy Score : 0.9764001309535905

AUC Score : 0.8924669134548138

Confusion Matrix :

```
[[170862  179]
 [ 4002  2119]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	1.00	0.99	171041
True	0.92	0.35	0.50	6121

# Hyperparameter Tuning

```
RandomForestClassifier()
```

```
(n_estimators: int = 100, *, criterion: str = "gini",  
max_depth: Any | None = None, min_samples_split: int =  
2, min_samples_leaf: int = 1, min_weight_fraction_leaf:  
float = 0, max_features: str = "sqrt", max_leaf_nodes:  
Any | None = None, min_impurity_decrease: float = 0,  
bootstrap: bool = True, oob_score: bool = False,  
n_jobs: Any | None = None, random_state: Any | None =  
None, verbose: int = 0, warm_start: bool = False,  
class_weight: Any | None = None, ccp_alpha: float = 0,  
max_samples: Any | None = None) -> None
```

```
# Parameters of Interest
```

```
criterion = ["gini", "entropy", "log_loss"]
```

```
bootstrap = [True, False]
```

```
# Grid Search
```

```
grid = dict(criterion=criterion, bootstrap=bootstrap)
```

```
grid_search = GridSearchCV(estimator=rfor, param_grid=grid, n_jobs=-1, cv=2, scoring='roc_auc')
```

```
grid_result = grid_search.fit(x_train, y_train)
```

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

```
Best: 0.909344 using {'bootstrap': False, 'criterion': 'entropy'}
```



# Results

## Original Random Forest

Accuracy Score : 0.9764001309535905

AUC Score : 0.8924669134548138

Confusion Matrix :

```
[[170862  179]
 [ 4002  2119]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	1.00	0.99	171041
True	0.92	0.35	0.50	6121

## Tuned Random Forest

Accuracy Score : 0.9777378896151545

AUC Score : 0.9138097188178325

Confusion Matrix :

```
[[170866  175]
 [ 3769  2352]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	1.00	0.99	171041
True	0.93	0.38	0.54	6121


AUC Score: 0.89 -> 0.91

F1 Score: 0.50 -> 0.54

# Thresholds

```
def to_labels(pos_probs, threshold):  
    return (pos_probs >= threshold)  
  
# Define thresholds  
thresholds = np.arange(0, 1, 0.01)  
# Evaluate each threshold  
scores = [f1_score(y_test, to_labels(y_prob_pred_rfor, t)) for t in thresholds]  
# Get the best threshold  
ix = np.argmax(np.array(scores))  
print(f"The best f1 score is {scores[ix]} and the threshold is {thresholds[ix]}")
```

The best f1 score is 0.611260568793236 and the threshold is 0.13



# Results

## Tuned Random Forest (w/o thresholds)

Accuracy Score : 0.9777378896151545

AUC Score : 0.9138097188178325

Confusion Matrix :

```
[[170866   175]
 [  3769  2352]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	1.00	0.99	171041
True	0.93	0.38	0.54	6121

AUC Score: 0.91 -> 0.90

## Tuned Random Forest (w/ thresholds)

Accuracy Score : 0.9771621453810636

AUC Score : 0.9071698020326076

Confusion Matrix :

```
[[169935   1106]
 [  2940   3181]]
```

Classification Report :

	precision	recall	f1-score	support
False	0.98	0.99	0.99	171041
True	0.74	0.52	0.61	6121

F-1 Score: 0.54 -> 0.61